

Linear Dynamic Systems Notes

Chris Beard

19 August 2011

Contents

1	Lecture 2	4
1.1	Interpretation for $y = Ax$	5
1.2	Example in notes	5
1.2.1	Linear elastic structure	5
1.2.2	Force/Torque on rigid body	5
1.2.3	Thermal System	5
2	Lecture 3	5
2.1	Review of Linearization (affine approximation)	5
2.1.1	Good intuition range-finding problem	6
2.2	Intepretations of $Ax = y$	6
2.2.1	Computing $C = AB$	6
2.2.2	Zero nullspace (in notes)	6
2.2.3	Interpretations of nullspace	6
2.2.4	Range	7
2.2.5	Rank of matrix	8
2.2.6	Various wrap-up items	9
3	Lecture 5	9
3.1	Geometric properties of orthonormal vectors	9
3.2	Orthonormal basis for \mathbb{R}^n	9
3.3	Expansion in orthonormal basis	10
3.4	Gram-Schmidt procedure	10
3.5	QR decomposition	10
3.6	General Gram Schmidt procedure ('rank revealing QR algorithm')	10
3.7	Applications	11
3.8	Least Squares Approximation	11
4	Lecture 6	11
4.1	Overdetermined equations	11
4.2	Least Squares 'Solution'	11
4.3	Projection on $\mathcal{R}(A)$	11
4.4	Orthogonality principle	11
4.5	Least-squares via QR factorization	12
4.6	Full QR factorization	12
4.7	Applications for least squares approximations	12
4.8	BLUE: Best linear unbiased estimator	12
4.9	Range-finding example	13
4.10	Quantizer example	13
4.11	Least Squares data-fitting	13
4.12	Least-squares polynomial fitting	13

5	Lecture 7	13
5.1	Least-squares polynomial fitting, cont'd	13
5.2	Growing sets of regressors	14
5.3	Least-squares system identification (important topic)	14
5.4	Model order selection	14
5.4.1	Cross-validation	15
5.5	Growing sets of measurements	15
5.6	Recursive ways to do least squares	15
5.7	Fast update algorithm for recursive LS	15
5.8	Multi-objective least squares	15
6	Lecture 8	16
6.1	Plot of achievable objective pairs	16
6.2	Minimizing weighted-sum objective	16
6.3	Example: frictionless table	16
6.4	Regularized least-squares	17
6.5	Nonlinear least squares (NLLS) problem	17
6.6	Gauss-Newton method for NLLS	18
6.7	G-N example	18
6.8	Underdetermined linear equations	18
6.9	Least norm solution	19
7	Lecture 9 pt 2	19
7.1	General norm minimization with equality constraints	19
7.2	Autonomous linear dynamical systems	19
7.3	Block diagrams	19
7.4	Linear circuit example	20
8	Lecture 10	20
8.1	Example: Series reaction $A \rightarrow B \rightarrow C$	20
8.2	Discrete time Markov chain	20
8.3	Numerical integration of continuous system	20
8.4	Higher order linear dynamical systems ($\dot{x} = Ax$)	20
8.5	Example: Mechanical systems	20
8.6	Linearization near equilibrium point	21
8.7	Example: pendulum linearization	21
9	Lecture 11	21
9.1	Laplace transform	21
9.2	Example: Harmonic oscillator	22
9.3	Example: Double Integrator	22
9.4	Characteristic polynomial	22
9.5	Get eigenvalues of A and poles of resolvent	22
9.6	Matrix exponential	22
9.7	Time transfer property	23
9.8	Application: sampling a continuous time system	23
10	Lecture 12	23
10.1	Stability	23
10.2	Eigenvectors and diagonalization	23
10.3	Scaling interpretation	23
10.4	Dynamic interpretation	24
10.5	Invariant set	24
10.6	Complex eigenvectors	24
10.7	Dynamic interpretation: left eigenvectors	24

10.8 Summary:	24
10.9 Example- companion matrix	24
11 Lecture 13	25
11.1 Example: Markov chain	25
11.2 Diagonalization	25
11.3 Not all matrices diagonalizable	25
11.4 Distinct eigenvalues	25
11.5 Diagonalization and left eigenvectors	26
11.6 Modal form	26
11.7 Real modal form	26
11.8 Diagonalization simplification	26
11.9 Simplify for analytical functions of a matrix	27
11.10 Solution via diagonalization	27
11.11 Interpretation	27
11.12 Application	27
11.13 Stability of discrete-time systems	27
11.14 Jordan Canonical form	27
11.15 Jordan block example	28
12 Lecture 14	28
12.1 J-C application: Caley-Hamilton theorem	28
12.2 Solving huge matrix equations	29
12.3 LTI system inputs & outputs	29
12.4 Interpretations	29
12.5 Transfer matrix	29
12.6 Impulse matrix	30
12.7 Step response	30
12.8 Spring dashpot example	30
12.9 RC circuit example	30
13 Lecture 15	30
13.1 DC or static gain matrix	30
13.2 Discretization with piecewise constant inputs	30
13.3 Causality	30
13.4 Idea of state	30
13.5 Change of coordinates	31
13.6 Standard forms for LDS	31
13.7 Discrete-time systems	31
13.8 Z-transform	31
13.9 Discrete-time transfer function	31
14 Last Section, Singular Value Decomposition	31
14.1 Eigenvalues of symmetric matrices	31
14.2 Eigenvectors of symmetric matrices	32
14.3 Interpretations	32
14.4 Example: RC circuit	33
15 Lecture 16	33
15.1 Quadratic forms	33
15.2 Examples of quadratic forms	33
15.3 Inequalities for quadratic forms	34
15.4 Ellipsoids	34
15.5 Gain of matrix in a direction	35
15.6 Example	35

15.7 Properties of matrix norm	35
16 Lecture 17	35
16.1 SVD	35
16.2 Dyadic expansion: (U and V have orthonormal columns)	36
16.3 Interpretations	36
16.4 General pseudo-inverse	36
16.5 Pseudo-inverse via regularization	37
17 Lecture 18	37
17.1 Sensitivity of linear equations to data error	37
17.2 Low rank applications	37
17.3 State transfer	37
17.4 Reachability	37
17.5 Reachability for discrete-time LDS	38
18 Lecture 19	38
18.1 General state transfer	38
18.2 Minimum time control problem	38
18.3 Least-norm input for reachability	39
18.4 Stability	39
18.5 Continuous time reachability	39
18.6 Impulsive inputs	39
18.7 Example	39
18.8 Least-norm input for reachability	39
19 Lecture 20	39
19.1 Observability and state estimation	39
19.2 Noiseless case	40
20 Appendix	40
20.1 Inverse, transpose properties	40
20.2 Invertibility implications	40
20.3 Tips	41
20.4 Unrelated stuff	41
20.4.1 Series	41
20.4.2 Taylor/Mclairen	42
21 Homework assignments	42

Following notes based on the EE263 Course Reader. Orgmode \rightarrow L^AT_EX 2_ε

Matlab files Linear Algebra
Index

1 Lecture 2

Official Lecture Notes

1.1 Interpretation for $y = Ax$

- A a transformation matrix
- y an observed measurement, x an unknown
- x is input, y is output; for rows i and columns j in A , $\mathbf{i} \rightarrow \mathbf{y}$, $\mathbf{j} \rightarrow \mathbf{x}$
 - indexed as **output, input**
- Lower diagonal Matrix should make you expect or have a vague thought about causality.
 - $a_{ij} = 0$ for $i < j \Rightarrow y_i$ only depends on x_1, \dots, x_i
 - A is diagonal; output only depends on input
- Sparsity pattern (block of zeroes) in a matrix should have you wonder why... usually not coincidental
- A_{35} positive $\Rightarrow x_5$ increased corresponds to an increase in 3rd output, y_3

1.2 Example in notes

1.2.1 Linear elastic structure

- a_{11} probably positive
 - x_1 input gives positive push to y_1 output

1.2.2 Force/Torque on rigid body

- Net torque/force on body is linearly related to force inputs

1.2.3 Thermal System

- despite normally needing a Poisson equation, the steady state heat of the different locations can be represented as a linear system $Ax = y$

2 Lecture 3

2.1 Review of Linearization (affine approximation)

1. If $f : \mathbf{R}^n \rightarrow \mathbf{R}^m$ is differentiable at $x_0 \in \mathbf{R}^n$, then x near $x_0 \Rightarrow f(x)$ very near $f(x_0) + Df(x_0)(x - x_0)$ where

$$Df(x_0)_{ij} = \left. \frac{\partial f_i}{\partial x_j} \right|_{x_0}$$

is the derivative (Jacobian) matrix.

2. with $y = f(x)$, $y_0 = f(x_0)$, define ‘input deviation’ $\delta x := x - x_0$, ‘output deviation’ $\delta y := y - y_0$
3. then we have $\delta y \approx Df(x_0)\delta x$
 - i.e., we get a linear function for looking at how the output changes with small changes in the input
4. When deviations are small, they are approximately related by a linear function

2.1.1 Good intuition range-finding problem**2.2 Interpretations of $Ax = y$**

- Scaled combination of columns in A

$$A = [a_1 a_2 \dots a_n] \rightarrow y = x_1 a_1 + x_2 a_2 + \dots + x_n a_n$$

- Looking at the rows: the multiplication is the inner product of rows and vector x
- I/O ordering is backwards in control; in A_{ij} , j refers to input and i refers to output
 - indexing is likewise backwards in block diagram indexing

$$AB = I; \tilde{a}_i^T \cdot b_j = 0 \text{ if } i \neq j, \text{ where } \tilde{a}_i \text{ is the } i\text{th row in } A, \text{ and } b_j \text{ is } j\text{th column in } B$$

2.2.1 Computing $C = AB$

How to compute this faster than using formula $c_{ij} = \sum_{k=1} A_{ik} B_{kj}$: have computer break it up into submatrices and do the multiplication on them

2.2.2 Zero nullspace (in notes)

- if 0 is only element in $\mathcal{N}(A)$
 - A is one-to-one
 - * linear transformation doesn't lose information
 - columns of A are independent, and basis for their span
 - A has a left inverse \Rightarrow you can use this to undo the transformation and find the input x , given the output y
 - $|A^T A| \neq 0$

2.2.3 Interpretations of nullspace

supposing $z \in \mathcal{N}(A)$

- Measurements
 - z is undetectable from sensors
 - x and $x + z$ are indistinguishable from sensors: $Ax = A(x + z)$
 - the nullspace characterizes ambiguity in x from measurement $y = Ax$
 - * large nullspace is bad
- $y = Ax$ is output from input x
 - z is input with no result
 - x and $x + z$ have same result
 - the nullspace characterizes freedom of input choice for given result
 - * large nullspace is good: more room for optimization because of more input possibilities

2.2.4 Range

Range of $A \in \mathbb{R}^{m \times n}$ defined as $\mathcal{R}(A) = \{Ax | x \in \mathbb{R}^n\} \subseteq \mathbb{R}^m$

- The set of vectors that can be ‘hit’ by linear mapping $y = Ax$
- span of columns of A
- set of vectors y for which $Ax = y$ has a solution
- possible sensor measurement
 - in design, you’ll want to throw an exception if a measurement is outside the range; the sensor is bad
 - test for a bad 13th sensor: remove 13th row in A ; if the reduced y is in the range of the reduced matrix, the 13th sensor might not have been bad
- Onto matrices

A is ‘onto’ if $\mathcal{R}(A) = \mathbb{R}^m$

 - you can solve $Ax = y$ for any y
 - columns of A span \mathbb{R}^m
 - A has a right inverse B s.t. $AB = I$
 - * can do $AB y = A(B y) = y$: you want an x that gives you y ? Here it is.
 - * Design procedure
 - rows of A are independent
 - * a.k.a., $\mathcal{N}(A^T) = \{0\}$
 - $|AA^T| \neq 0$
- Interpretations of range
 - supposing $v \in \mathcal{R}(A)$
 - * v reachable
 - * else, not reachable
- Inverse

Note: square matrices are impractical for engineering. They don’t let you take advantage of redundant sensors/controllers, or let you build a robust system to take care of broken sensors

 - $A \in \mathbb{R}^{n \times n}$ is invertible or nonsingular if $\det A \neq 0$
 - * columns of A are basis for \mathbb{R}^n
 - * rows of A are basis for \mathbb{R}^n
 - * $y = Ax$ has a unique solution x for every $y \in \mathbb{R}^n$
 - * A has left and right inverse $A^{-1} \in \mathbb{R}^{n \times n}$, s.t. $AA^{-1} = A^{-1}A = I$
 - * $\mathcal{N}(A) = \{0\}$
 - * $\mathcal{R}(A) = \mathbb{R}^n$
 - * $\det A^T A = |AA^T| \neq 0$
 - Dual basis interpretation of inverse

a_i are columns of A , and \tilde{b}_i^T are rows of $B = A^{-1}$

 - * $y = Ax$, column by column, looks like $y = x_1 a_1 + \dots + x_n a_n$
 - multiply both sides of $y = Ax$ by $A^{-1} = B$ gives $x = By$
 - so $x_i = \tilde{b}_i^T y$

$$\begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ a_1 & a_2 & \dots & a_n \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$x = A^{-1}y$$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} \dots & \tilde{b}_1^T & \dots \\ \dots & \tilde{b}_2^T & \dots \\ \dots & \vdots & \dots \\ \dots & \tilde{b}_n^T & \dots \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = [x_1 a_1 + \dots + x_n a_n] = [(\tilde{b}_1^T y) a_1 + \dots + (\tilde{b}_n^T y) a_n]$$

Beautiful thing:

$$y = \sum_{i=1}^n (\tilde{b}_i^T y) a_i$$

2.2.5 Rank of matrix

Rank of $A \in \mathbb{R}^{m \times n}$ as $\mathbf{rank}(A) = \mathbf{dim} \mathcal{R}(A)$

- $\mathbf{rank}(A) = \mathbf{rank}(A^T)$
- $\mathbf{rank}(A)$ is maximum number of independent columns or rows of A : $\mathbf{rank}(A) \leq \mathbf{min}(m, n)$
- $\mathbf{rank}(A) + \mathbf{dim} \mathcal{N}(A) = n$
- Conservation of degrees of freedom (dimension)
 - $\mathbf{rank}(A)$ is dimension of set ‘hit’ by mapping $y = Ax$
 - $\mathbf{dim} \mathcal{N}(A)$ is dimension of set of x ‘crushed’ to zero by $y = Ax$
 - Example
 - * $A \in \mathbb{R}^{20 \times 10}$ $\mathbf{rank}(A) = 8$
 - you can do 8 dimensions worth of stuff
 - 10 knobs, 2 redundant knobs, which is $\mathbf{dim} \mathcal{N}(A) = 2$
- Coding interpretation of rank
 - rank of product: $\mathbf{rank}(BC) \leq \mathbf{min}\{\mathbf{rank}(B), \mathbf{rank}(C)\}$
 - supposedly really cool stuff based on this
 - low rank matrices let you do fast computations

2.2.6 Various wrap-up items

- RMS

$$\text{rms}(x) = \left(\frac{1}{n} \sum_{i=1}^n \right)^{1/2} = \frac{\|x\|}{\sqrt{n}}$$

- Inner product

$$\langle x, y \rangle := x_1 y_1 + x_2 y_2 + \cdots + x_n y_n = x^T y$$

- interpretation of inner product signs:
 - $x^T y > 0$: acute; roughly point in same direction
 - $x^T y < 0$: obtuse; roughly point in opposite direction

- Orthonormal set of vectors

- set of k vectors $u_1, u_2, \dots, u_k \in \mathbb{R}^n$ orthonormal; $U = [u_1 \cdots u_k]$
- $U^T U = I_k \Leftrightarrow$ set of column vectors of U are orthonormal
- **warning:** $U U^T \neq I_k$ if $k < n$
 - * say U is 10×3 , U^T is 3×10 , rank of U is 3 \Rightarrow rank of $U U^T$ is at most 3
 - * but $U U^T$ will be a 10×10 matrix, so it can't be the identity matrix

3 Lecture 5

A good source for more on orthogonality at University of Minnesota

3.1 Geometric properties of orthonormal vectors

- columns of U are ON \Rightarrow mapping under U preserves distances

$$w = Uz \Rightarrow \|w\| = \|z\|$$

- Also preserves inner product
- Also preserves angles
- Something like a rigid transformation

3.2 Orthonormal basis for \mathbb{R}^n

- if there are n orthonormal vectors (remember, with dimension n), it forms an orthonormal basis for \mathbb{R}^n
- $U^{-1} = U^T$
 - $U^T U = I \Leftrightarrow U$'s column vectors form an orthonormal basis for \mathbb{R}^n
 - $\sum_{i=1}^n u_i u_i^T = I \in \mathbb{R}^{n \times n}$ (known as a dyad, or outer product; inner products reverses the two and gives a scalar, outer gives a matrix)
 - outer products take 2 vectors, possibly of different sizes, and multiplies every combination of elements one with another

3.3 Expansion in orthonormal basis

- U orthogonal $\Rightarrow x = UU^T$
- $x = \sum_{i=1}^n (u_i^T x) u_i$
 - because $U^T U = I$, the thing in sum is really $u_i u_i^T x$
 - $u_i^T x$ is really a scalar, so this can be moved to the front of u_i , giving our result
 - This says x is a linear combination of u_i 's

3.4 Gram-Schmidt procedure

- $a_1, \dots, a_k \in \mathbb{R}^n$ are LI; G-S finds ON vectors q_1, \dots, q_k s.t.

$$\text{span}(a_1, \dots, a_r) = \text{span}(q_1, \dots, q_r)$$

for $r \leq k$

- so q_1, \dots, q_r is an ON basis for $\text{span}(a_1, \dots, a_r)$
- Basic method: orthogonalize each vector wrt the previous ones, then normalize result
 1. $\tilde{q}_1 = a_1$
 2. normalize: $q_1 = \tilde{q}_1 / \|\tilde{q}_1\|$
 3. remove q_1 component from a_2 : $\tilde{q}_2 = a_2 - (q_1^T a_2) q_1$
 4. normalize q_2
 5. remove q_1, q_2 components: $\tilde{q}_3 = a_3 - (q_1^T a_3) q_1 - (q_2^T a_3) q_2$
 6. normalize q_3
- $a_i = (q_1^T a_i) q_1 + (q_2^T a_i) q_2 + \dots + (q_{i-1}^T a_i) q_{i-1} + \|\tilde{q}_i\| q_i$
 - $= r_{1i} q_1 + r_{2i} q_2 + \dots + r_{ii} q_i$ ($r_{ii} \geq 0$ is the length of \tilde{q}_i)

3.5 QR decomposition

This can be written as $A = QR$, where $A \in \mathbb{R}^{n \times k}$, $Q \in \mathbb{R}^{n \times k}$, $R \in \mathbb{R}^{k \times k}$

$$\begin{bmatrix} a_1 & a_2 & \cdots & a_k \end{bmatrix} = \begin{bmatrix} q_1 & q_2 & \cdots & q_k \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1k} \\ 0 & r_{22} & \cdots & r_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & r_{kk} \end{bmatrix}$$

- R triangular because computation of a_i only involves up to q_i
 - a sort of causality, since you can calculate q_7 without seeing q_8
- Columns of Q are ON basis for $\mathcal{R}(A)$

3.6 General Gram Schmidt procedure ('rank revealing QR algorithm')

- Basically the same, but if one of the \tilde{q}_i 's is zero (meaning a_i is dependent on previous a vectors), then just go to the next column
- referring to notes, upper staircase notation shows which vectors are dependent on previous ones (columns without the x's)
 - entries with x are 'corner' entries

3.7 Applications

- check if $b \in \text{span}(a_1, a_2, \dots, a_k)$
- Factorize matrix A

3.8 Least Squares Approximation

- Overdetermined linear equation (tall, skinny, more equations than unknowns, dimensionally redundant system of equations)

4 Lecture 6

On skinny, full rank matrices

4.1 Overdetermined equations

- Skinny, more equations than unknowns
- Given $y = Ax$, $A \in \mathbb{R}^{m \times n}$, a randomly-chosen y in \mathbb{R}^m has 0 probability of being in the range of A
- To *approximately* solve for y , minimize norm of error (residual) $r = Ax - y$
- find $x = x_{ls}$ (least squares approx.) that minimizes $\|r\|$

4.2 Least Squares ‘Solution’

- square $\|r\|$, get expansion, set gradient wrt x equal to zero
- $x_{ls} = (A^T A)^{-1} A^T y = B_{ls} y$ (linear operation)
- $A^T A$ should be invertible, square, full rank
- $(A^T A)^{-1} A^T$ is a generalized inverse (is only inverse for square matrices, though)
 - Also known as the A^\dagger , ‘pseudo-inverse’
 - Which is a left inverse of A

4.3 Projection on $\mathcal{R}(A)$

Ax_{ls} is the point closest to y (i.e., projection of y onto $\mathcal{R}(A)$)

- $Ax_{ls} = \text{proj}_{\mathcal{R}(A)}(y) = (A(A^T A)^{-1} A^T) y$

4.4 Orthogonality principle

The optimal residual is orthogonal to $C(A)$

- $r = Ax_{ls} - y = (A(A^T A)^{-1} A^T - I)y$ orthogonal to $C(A)$
- $\langle r, Az \rangle = y^T (A(A^T A)^{-1} A^T - I)^T A z = 0$ for all $z \in \mathbb{R}^n$

4.5 Least-squares via QR factorization

A is still skinny, full rank

- Factor as $A = QR$; $Q^T Q = I_n$, $R \in \mathbb{R}^{n \times n}$ upper triangular, invertible
- pseudo-inverse: $(A^T A)^{-1} A^T = R^{-1} Q^T \Rightarrow \boxed{R^{-1} Q^T y = x_{ls}}$
- Pretty straight-forward
- Matlab for least squares approximation

```
x1 = inv(A' * A)*A'y; # So common that has shorthand in MATLAB
x1 = A\y;             # Works for non-skinny matrices, may do unexpected things
```

4.6 Full QR factorization

- $A = \begin{bmatrix} Q_1 & Q_2 \end{bmatrix} \begin{bmatrix} R_1 \\ 0 \end{bmatrix}$
 - New Q is square, orthogonal matrix; R_1 is square, upper triangular, invertible
- Remember, multiplying by orthogonal matrix doesn't change the norm:
 - $\|Ax - y\|^2 = \|R_1 x - Q_1^T y\|^2 + \|Q_2^T y\|^2$
 - Find least squares approximation with $x_{ls} = R_1^{-1} Q_1^T y$ (zeroes first term)

4.7 Applications for least squares approximations

- if there is some noise v in $y = Ax + v$
 - you can't reconstruct x , but you can get close with the approximation
- Estimation: choose some \hat{x} that minimizes $\|A\hat{x} - y\|$, which is the deviation between the think we observed, and what we would have observed in the absence of noise

4.8 BLUE: Best linear unbiased estimator

- A still full rank and skinny; have a 'linear estimator' $\hat{x} = By$ (B is fat)
 - $\hat{x} = B(Ax + v)$
- Called unbiased if there is no estimation error when there's no noise; the estimator works perfectly in the absence of noise
 - if $v = 0$ and $BA = I$; B is left inverse/perfect reconstructor
- Estimation error of unbiased linear estimator is $x - \hat{x} = sBv$, so we want B to be small and $BA = I$; small means error isn't sensitive to the noise
- The pseudo-inverse is the smallest left inverse of A :

$$\begin{aligned}
 & - A^\dagger = (A^T A)^{-1} A^T \\
 & - \sum_{i,j} B_{ij}^2 \geq \sum_{i,j} A_{ij}^{\dagger 2}
 \end{aligned}$$

4.9 Range-finding example

- Find ranges to 4 beacons from an unknown position x
- $y = - \begin{bmatrix} k_1^T \\ k_2^T \\ k_3^T \\ k_4^T \end{bmatrix} x + v$
- actual position $x = (5.59, 10.58)$; measurement $y = (-11.95, -2.84, -9.81, 2.81)$
 - these numbers aren't consistent in $Ax = y$, since there's also the error; there is no such x value that can give this y value
- There are 2 redundant sensors (2 more y values than x values); one method for estimating \hat{x} is 'just enough' method: you only need 2 y values; take inverse of top half of A and pad the rest of the matrix with 0's
- use $\hat{x} = B_{justenough}y = \begin{bmatrix} [k_1^T]^{-1} & 0 & 0 \\ k_2^T & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0 & -1.0 & 0 & 0 \\ -1.12 & .5 & 0 & 0 \end{bmatrix} y = \begin{bmatrix} 2.84 \\ 11.9 \end{bmatrix}$
- Least Squares method: $\hat{x}A^\dagger y$ = this has a much smaller norm of error
- Just enough estimator doesn't seem to have good performance... unless last two measurements were really off, since JEM only takes 2 measurements into account

4.10 Quantizer example

Super-impressive least squares estimate; more precise than A-D converter

4.11 Least Squares data-fitting

- use functions $f_1, f_2, \dots, f_n : S \rightarrow \mathbb{R}$ are called regressors or basis functions
- applications
 - interpolation-, extrapolation, smoothing of data

Applications

interpolation	don't have sensors in specific location, but want the temperature
extrapolation	get good basis functions for better interpolation
data smoothing	de-noise measurements
simple, approximate data model	Get a million samples, use the data-fitting to get a simple approximate function

4.12 Least-squares polynomial fitting

- Vandermonde matrix?

5 Lecture 7

5.1 Least-squares polynomial fitting, cont'd

- have data samples $(t_i, y_i), i = 1, \dots, m$
- fit coefficients a_i of polynomial $p(t) = a_0 + a_1t + \dots + a_{n-1}t^{n-1}$ so that when evaluated at t_i it will give you the associated y value
- basis functions are $f_j(t) = t^{j-1}, j = 1, \dots, n$

- use Vandermonde matrix A ('polynomial evaluator matrix'):

$$A = \begin{bmatrix} 1 & t_1 & t_1^2 & \dots & t_1^{n-1} \\ 1 & t_2 & t_2^2 & \dots & t_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & t_m & t_m^2 & \dots & t_m^{n-1} \end{bmatrix}$$

- side note: use this when you want to fit throughout an interval, use a Taylor series fit if you want it close to a point

5.2 Growing sets of regressors

- Given ordered set of vectors; find best fit with first vector, then best fit with first and second, then best fit with first three...
- These vectors called *regressors*, or columns
- Say you have some *master list* A with n columns, and $A^{(p)}$ will be the matrix with the first p columns of it
 - we want to minimize different sets of $\|A^{(p)}x - y\|$
 - i.e., project y onto a growing span $\{a_1, a_2, \dots, a_p\}$
- Solution for each $p \leq n$ given by $x_{ls}^{(p)} = (A_p^T A_p)^{-1} A_p^T y = R_p^{-1} Q_p^T y$
 - In MATLAB, `A(:,1:p)\y`, though technically it's faster to do a sort of **for** loop
- Residual, $\|\sum_{i=1}^p x_i a_i - y\|$ reduces as p (number of columns) increases
 - though it may be same as residual with previous value of p if the optimal $x_1 = 0$, when $y \perp a_1$
 - if the residual drops 15% from that of previous value of p , you say that a_1 explains 15% of y

5.3 Least-squares system identification (important topic)

- measure input, output $u(t), y(t)$ for $t = 0, \dots, N$ of unknown system, and try to get a model of system
- example: moving average (MA) model with n delays (try to approximate what are the weights h_i for each delay)
 - see equation/matrix in notes, though there are different ways to write it
 - get best answer with LSA

5.4 Model order selection

- how large should n be?
- the larger, the smaller prediction error on *data used to form model*
- but at a certain point, predictive ability of model on other I/O data from same system worsens
- probably best to choose the 'knee' on the graph on notes slide for prediction of new data

5.4.1 Cross-validation

- check with new data, only if you're getting small residuals on data you've already seen
- when n gets too large (greater than $n = 10$ on graph), the error with 'validation data' actually gets larger
- this example is ideal, since $n = 10$ is the obvious order for the model
- **Application note:** in medical, many industries, there's a firm wall between validation data and model-developing data, so someone *else* tests your model
- in this example, it is known as *overfit* when the validation data error gets larger for n too large

5.5 Growing sets of measurements

- similar to GSo Regressors, except you add new rows, not columns
- this would happen if we're estimating a parameter x (which is constant)
- Solution: $x_{ls} = \left(\sum_{i=1}^m a_i a_i^T \right)^{-1} \sum_{i=1}^m y_i a_i$
- new way to think of least squares equation

5.6 Recursive ways to do least squares

- don't have to re-add for each new measurement
 - i.e., memory is bounded
 - use equation from notes; solution is $x_{ls}(m) = P(m)^{-1}q(m)$

5.7 Fast update algorithm for recursive LS

- Was a big deal back in the day; somewhat still

5.8 Multi-objective least squares

- Sometimes you have 2+ objectives to minimize
 - say $J_1 = \|Ax - y\|^2$ (what we've done so far)
 - and $J_2 = \|Fx - g\|^2$
 - these are usually competing (minimize one at cost of other)
- Variable in question is $x \in \mathbb{R}^n$
- Plot in notes shows plot of $(J_1(x_i), J_2(x_i))$
- Some points are unambiguously worse than others, but there is some ambiguity when $J_1(x_1) < J_1(x_2)$, while $J_2(x_1) > J_2(x_2)$
- Fix this ambiguity with 'weighted-sum objective'
- $J_1 + \mu J_2 = \|Ax - y\|^2 + \mu \|Fx - g\|^2$
 - Say, there's a trade-off between smoothness (no noise) and better fit; μ can have different dimensions if J_2 does
- Use slope of μ in graph ('indifference curve', in economics) [slide 7-6]

6 Lecture 8

Multi-objective least-squares

6.1 Plot of achievable objective pairs

- if it approximates an L shape (has a 'knee'), the knee is usually the obvious optimal location, so least-squares isn't as helpful
 - optimal point isn't very sensitive to μ
- Other extreme: trade-off curve looks linear (negative slope), where it's zero-sum
 - optimal point very sensitive to μ
 - slope commonly called *exchange rate curve*
- In this class, they must be convex curves (cup up/outward)
- To find Pareto optimal points, minimize $J_1 + \mu J_2 = \alpha$
 - on plot, can have level curves with slope μ
 - Find point on Pareto Optimal Curve that has slope μ

6.2 Minimizing weighted-sum objective

- note: norm-squared of a stacked vector is norm-square of the top+norm-square of bottom

$$J_1 + \mu J_2 = \|Ax - y\|^2 + \mu \|Fx - g\|^2 = \left\| \begin{bmatrix} A \\ \sqrt{\mu}F \end{bmatrix} x - \begin{bmatrix} y \\ \sqrt{\mu}g \end{bmatrix} \right\|^2$$

$$= \left\| \tilde{A}x - \tilde{y} \right\|^2$$

where

$$\tilde{A} = \begin{bmatrix} A \\ \sqrt{\mu}F \end{bmatrix}, \tilde{y} = \begin{bmatrix} y \\ \sqrt{\mu}g \end{bmatrix}$$

If \tilde{A} is full rank,

$$x = \left(\tilde{A}^T \tilde{A} \right)^{-1} \tilde{A}^T \tilde{y} \tag{1}$$

$$= (A^T A + \mu F^T F)^{-1} (A^T y + \mu F^T g) \tag{2}$$

Note: to plot the tradeoff curve, calculate the minimizer x_μ , and plot the resulting pairs (J_1, J_2) In MATLAB, `[A; sqrt(mu) * F] \ [y; sqrt(mu) * g]`

6.3 Example: frictionless table

- y is final position at $t = 10$; $y = a^T x$, $a \in \mathbb{R}^{10}$
- $J_1 = (y - 1)^2$, (final position difference from $y = 1$ squared)
- $J_2 = \|x\|^2$ sum of force squares
- Q: Why do we often care about sum of squares? A: **It's easy to analyze** (not necessarily because it corresponds to energy)
 - $\max |x_i|$ corresponds to maximum thrust
 - $\sum |x_i|$ corresponds to fuel use
- Optimal tradeoff curve is quadratic

6.4 Regularized least-squares

- famous example of multi-objective least squares

– second J term is simply $J_2 = \|x\|$, though first is the same: $J_1 = \|Ax - y\|^2$

- Tychonov regularization works for *any* A

– *regularized* least-squares solution: $x_\mu = (A^T A + \mu I)^{-1} A^T y$, for $F = I, g = 0$

Show $(A^T A + \mu I)$ is invertible, no matter what size/values of A (assuming $\mu > 0$): If this is *not* invertible (singular), it means some nonzero vector z gets mapped to zero ($z \in \mathcal{N}(A)$)

$$(A^T A + \mu I)z = 0, z \neq 0 \quad (3)$$

$$z^T (A^T A + \mu I)z = 0 \text{ since } z^T \vec{0} = 0 \quad (4)$$

$$z^T A^T A z + \mu z^T z = 0 \quad (5)$$

$$\|Az\|^2 + \mu \|z\|^2 = 0 \quad (6)$$

$$z = \vec{0} \quad (7)$$

So, z can only be zero, meaning $\mathcal{N}(A) = \{0\} \Rightarrow (A^T A + \mu I)$ is invertible. This is also why μ must be positive. Or, you know it's invertible, since it is full rank (and skinny) when you stack μI below it (see definition of \tilde{A}).

- Application of Regularized least-squares

- estimation/inversion
- $Ax - y$ is sensor residual
- prior information that x is really small
- or, model only accurate for small x
- Tychonov solution trades off sensor fit and size of x

- Image processing example

- Laplacian regularization
 - * image reconstruction problem
- x is vectorized version of image
- $\|Ax - y\|^2$ is difference from real image
- Want new objective to minimize roughness
 - * vector Dx (from new matrix D) which has difference between neighboring pixels as elements
 - $D_v x$ measures vertical difference
 - $D_h x$ measures horizontal difference
 - Nullspace is vector where there is no variation between pixels
- minimize $\|Ax - y\|^2 + \mu \| [D_h x \ D_v x]^T \|^2$
 - * if μ is turned way up, it'll be all smoothed out
 - * if you care about total size of image, you can add another parameter λ : $\|Ax - y\|^2 + \mu \| [D_h x \ D_v x]^T \|^2 + \lambda \|x\|^2$

6.5 Nonlinear least squares (NLLS) problem

- find $x \in \mathbb{R}^n$ that minimizes $\|r(x)\|^2 = \sum_{i=1}^m r_i(x)^2$
- $r(x)$ is vector of residuals; $r(x) = Ax - y \Rightarrow$ problem reduces to linear least squares problem
- in general, can't **really** solve a NLLS problem, but can find good heuristics to get a locally optimal solution

6.6 Gauss-Newton method for NLLS

- Start guess for x
- Loop
 - linearize r near current guess
 - new guess is linear LS solution, using linearized r
 - if convergence, stop
- Linearize?
 - Jacobian: $(Dr)_{ij} = \partial r_i / \partial x_j$
 - Linearization: $r(x) \approx r(x^{(k)}) + Dr(x^{(k)})(x - x^{(k)})$
 - Set this linearized approximation equal to $r(x) \approx A^{(k)}x - b^{(k)}$
 - * $A^{(k)} = Dr(x^{(k)})$
 - * $b^{(k)} = Dr(x^{(k)})x^{(k)} - r(x^{(k)})$
 - See rest in notes
 - At k th iteration, approximate NLLS problem by linear LS problem:
 - * $\|r(x)\|^2 \approx \|A^{(k)}x - b^{(k)}\|^2$
 - if you wanna make this really cool add a $\mu\|x - x^{(k)}\|^2$ term on RHS
 - called a ‘trust region term’;
 - first (original) part says to minimize sum of squares for *model*
 - trust region term says ‘but don’t go far from where you are now’
- Could also linearize without calculus; works really well
 - See ‘particle filter’

6.7 G-N example

- Nice graph and residual plot
- As practical matter, good to run simulation several times (with different initial guesses)
- ‘exhaustive simulation’

6.8 Underdetermined linear equations

- $A \in \mathbb{R}^{m \times n}$, $m < n$ (A is fat)
- more variables than equations
- x is underspecified
- For this section **assume A is full rank**
- Set of all solutions has form $\{x | Ax = y\} = \{x_p + z | z \in \mathcal{N}(A)\}$
- solution has $\dim \mathcal{N}(A) = n - m$ ‘degrees of freedom’
 - many DOF: good for design (flexibility), bad for estimation (stuff you don’t/can’t know with available measurements)

6.9 Least norm solution

- $x_{ls} = A^T(AA^T)^{-1}y$
 - similar to our familiar skinny A version: $x_{ls} = (A^T A)^{-1}A^T y$
 - mnemonic: $(\cdot)^{-1}$ thing must be square
 - * if A skinny, both $A A^T$ and $A^T A$ could be square (syntactically)
 - * semantically, you need the up and down patterns that will form the smallest square, i.e., full rank matrix

7 Lecture 9 pt 2

Thank you fucking Suncheon.

7.1 General norm minimization with equality constraints

- Problem: minimize $\|Ax - b\|$ subject to $Cx = d$, with variable x
- Least squares/least norm are special cases
 - Least norm: set $A = I, b = 0$, then you just have norm of x subject to some linear equations
- Same as: minimize $(1/2)\|Ax - b\|^2$ subject to $Cx = d$
- Lagrangian is... long ugly thing... look at notes
 - a bit easier to look at block matrix format

$$\begin{bmatrix} A^T A & C^T \\ C & 0 \end{bmatrix} \begin{bmatrix} x \\ \lambda \end{bmatrix} = \begin{bmatrix} A^T b \\ d \end{bmatrix}$$

- recover least squares (maybe) by eliminating C from matrix (not setting to zero, but only having 1 row/column in first matrix)

7.2 Autonomous linear dynamical systems

“What the class is nominally about”

- In continuous time, autonomous LDS has form $\dot{x} = Ax$
- Solution: $x(t) = e^{ta}x(0)$
- $x(t) \in \mathbb{R}^n$ is called the state
 - n is state dimension
- A basically maps where you are (x) to where you’re going (\dot{x})
 - has units of s^{-1} , frequency
- Example illustration: vector fields

7.3 Block diagrams

- use integrators to express $\dot{x} = Ax$ instead of differentiators
 - block called ‘bank of integrators’
 - historically used because of analog, mechanical computers
- notches to express n signals

7.4 Linear circuit example

8 Lecture 10

Examples of autonomous linear dynamical systems, $\dot{x} = Ax$

8.1 Example: Series reaction $A \rightarrow B \rightarrow C$

$$\dot{x} = \begin{bmatrix} -k_1 & 0 & 0 \\ k_1 & -k_2 & 0 \\ 0 & k_2 & 0 \end{bmatrix} x$$

- For second row, first term on rhs of $\dot{x}_2 = k_1x_1 - k_2x_2$ is *buildup*
- Note: Column sums are 0 implies conservation of mass/materials;

8.2 Discrete time Markov chain

- $x(t+1) = Ax(t)$
- $x(t) = A^t x(0)$
- Given current state, the matrix of *transition probabilities* P will tell you probabilities of the next state, given the current state

8.3 Numerical integration of continuous system

- for a small time step h , find about where you'll be in h seconds
- $x(t+h) \approx x(t) + h\dot{x}(t) = (I + hA)x(t)$
- problem: when you do it for a long time, error can build up pretty high

8.4 Higher order linear dynamical systems ($\dot{x} = Ax$)

$$x^{(k)} = A_{k-1}x^{(k-1)} + \dots + A_1x^{(1)} + A_0x, x(t) \in \mathbb{R}^n$$

- define new variable

$$z = \begin{bmatrix} x \\ x^{(1)} \\ \vdots \\ x^{(k-1)} \end{bmatrix} \in \mathbb{R}^{nk}, \dot{z} = \begin{bmatrix} x^{(1)} \\ \vdots \\ x^{(k)} \end{bmatrix} = \begin{bmatrix} 0 & I & 0 & \cdots & 0 \\ 0 & 0 & I & \cdots & 0 \\ \vdots & & & & \vdots \\ 0 & 0 & 0 & \cdots & I \\ A_0 & A_1 & A_2 & \cdots & A_{k-1} \end{bmatrix} z$$

- 'upshift x , and zero-pad'
- z is the state, not x
- in notes, block diagram with chain of integrators

8.5 Example: Mechanical systems

- Ex: K_{12} is 'cross-stiffness', how much stiffness you'd feel at node 1 from node 2

8.6 Linearization near equilibrium point

Equilibrium point corresponds to constant solution ($f(x_e) = 0, x(t) = x_e$)

- if you start at an equilibrium point, you'll stay there
- if you start *near* equilibrium point
 - veer off (unstable)
 - go towards equilibrium (stable)
 - something in between
- but, you never stay at an unstable equilibrium position, since equation is really $\dot{x} = f(x) + w(t)$, where $w(t)$ is noise
- Near equilibrium point, $\dot{\delta x}(t) \approx Df(x_e)\delta x(t)$, where D is the Jacobian
 - similar to euler forward equation
- Don't fully trust approximations on approximations (but hope they work)

8.7 Example: pendulum linearization

- $ml^2\ddot{\theta} = -lmg \sin \theta$
- rewrite as 1st order DE with state $x = [\theta \ \dot{\theta}]^T = [x_1 \ x_2]^T$:

$$\dot{x} = \begin{bmatrix} x_2 \\ -(g/l) \sin x_1 \end{bmatrix}$$

- \exists equilibrium point at $x = 0$ (and π), so we linearize system near $x_e = 0$, using a Jacobian matrix:

$$\dot{\delta x} = \begin{bmatrix} \frac{\partial x_2}{\partial x_1} (-g/l) \sin x_1 |_{x_1=0} & \frac{\partial x_2}{\partial x_2} (-g/l) \sin x_1 \\ \frac{\partial}{\partial x_1} (-g/l) \sin x_1 |_{x_1=0} & \frac{\partial}{\partial x_2} (-g/l) \sin x_1 \end{bmatrix} \delta x = \begin{bmatrix} 0 & 1 \\ -g/l & 0 \end{bmatrix} \delta x$$

9 Lecture 11

Solution via Laplace transform and matrix exponential Remember, we've already overloaded $\dot{x} = ax$. Now, we'll overload exponentials to apply to matrices $x(t) = e^{ta}x(0)$.

9.1 Laplace transform

- $z : \mathbb{R}_+ \rightarrow \mathbb{R}^{p \times q}$ (function that maps non-negative real scalars to matrices)
- Laplace transform: $Z = \mathcal{L}(z)$, defined by $Z(s) = \int_0^\infty e^{-st} z(t) dt$
- Region of convergence of Z is mostly for confusing students
- Derivative property: $\mathcal{L}(\dot{z}) = sZ(s) - z(0)$

So, we can use the Laplace transform to solve $\dot{x} = Ax$. Take Laplace: $sX(s) - x(0) = AX(s)$, rewrite as $(sI - A)X(s) = x(0)$, so $X(s) = (sI - A)^{-1}x(0)$. Then take the inverse transform: $x(t) = \mathcal{L}^{-1}((sI - A)^{-1})x(0)$

- takes advantage if linearity of the Laplace transform
- $(sI - A)^{-1}$ is called the *resolvent* of A
 - but not defined for eigenvalues of A ; s , ST $\det(sI - A) = 0$
- $\Phi = \mathcal{L}^{-1}((sI - A)^{-1})$ is called the *state-transition matrix*, which maps the initial state to state at time t :
 $x(t) = \Phi(t)x(0)$

9.2 Example: Harmonic oscillator

$$\dot{x} = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} x$$

- To solve for s , get the resolvent, then apply the Laplacian to it *elementwise*, getting

$$x(t) = \begin{bmatrix} \cos t & \sin t \\ -\sin t & \cos t \end{bmatrix} x(0)$$

Which is a circular rotation matrix. The solutions to $\dot{x} = ax$ is $x(t) = e^{ta}x(0)$

- a positive: exponential growth
- a negative: exponential decay
- $a = 0$: constant

9.3 Example: Double Integrator

- Note, with scalars, x in $\dot{x} = ax$ grows exponentially in time, and cannot grow linearly, as with matrices (can have a t element in matrix)
- What is first column of $\Phi(t)$ say? It tells what the state trajectory is if the initial condition was e_1 (second column tells what it is if $x(0) = e_2$)
- First row says the linear combination that x_1 is at time t given $x(0)$

9.4 Characteristic polynomial

$\mathcal{X}(s) = \det(sI - A)$; called a *monic* polynomial

- roots of \mathcal{X} are eigenvalues of A , and \mathcal{X} has real coefficients, so e-values are real or occur in conjugate pairs

9.5 Get eigenvalues of A and poles of resolvent

Use Cramer's rule to get i, j entry:

$$(-1)^{i+j} \frac{\det \Delta_{ij}}{\det(sI - A)},$$

where Δ_{ij} is $sI - A$ with j th row and i th column deleted. Poles of entries of resolvent **must** be eigenvalues of A .

9.6 Matrix exponential

How to overload exponentials for matrices; start with $(I - C)^{-1} = I + C + C^2 + \dots$. Series converges if $|\text{eigenvalues of } C| < 1$. Do series expansion of resolvent, then take the Laplacian of the series, which looks like the form for the expansion of e^{ta} (though square matrices replace scalars). So we end by learning that the state transition matrix, $\Phi(t)$ is the matrix exponential e^{tA} .

- Many scalar exponential properties don't extend to matrix exponential; with scalars, this is wrong: $e^{A+B} = e^A e^B$ (unless A and B commute: $AB = BA$)
- But this is ok: $e^{-A} = (e^A)^{-1}$
- So, how do you find the matrix exponential:

Find e^A ,

$$A = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$$

Found $e^{tA} = \mathcal{L}^{-1}(sI - A)^{-1}$ in earlier example, so just plug in $t = 1$.

- Matlab: `expm(A)`, not elementwise `exp(A)`

9.7 Time transfer property

Summary: for $\dot{x} = Ax$, $x(t) = \Phi(t)x(0) = e^{tA}x(0)$. The matrix e^{tA} propagates initial condition into state at time t . Also propagates backward in time if $t < 0$.

If given $x(12)$, find $x(0)$ via $e^{-12A}x(12)$.

- Can use first order forward Euler approximate state update for small t
- Discretized autonomous LDS: $z(k+1) = e^{hA}z(k)$ (not an approximation for these equations)

9.8 Application: sampling a continuous time system

10 Lecture 12

Piecewise constant system: A is constant for certain intervals of time.

- Qualitative behavior of $x(t)$
 - Eigenvalues determine (possible) behavior of x
 - Can plot eigenvalues on complex axes; like pole plot
 - Can put x in summation form with polynomial coefficient and exponential terms

10.1 Stability

- $\dot{x} = Ax$ is stable if $e^{tA} \rightarrow 0$ as $t \rightarrow \infty$
 - means that state $x(t)$ converges to 0 as $t \rightarrow \infty$, no matter $x(0)$
 - all trajectories of $\dot{x} = Ax$ converge to 0 as $t \rightarrow \infty$
 - $\dot{x} = Ax$ is stable iff all eigenvalues of A have negative real part

10.2 Eigenvectors and diagonalization

- $\lambda \in \mathbb{C}$ is an eigenvalue of $A \in \mathbb{C}^{n \times n}$ if (characteristic polynomial)

$$\mathcal{X}(\lambda) = \det(\lambda I - A) = 0$$

- i.e., $(\lambda I - A)$ is singular, not invertible, \mathcal{N} not equal to the 0 set

Equivalent to:

- \exists nonzero $v \in \mathbb{C}^n$ s.t. $(\lambda I - A)v = 0$: $Av = \lambda v$ (v is the eigenvector)
 - columns are dependent
- \exists nonzero $w \in \mathbb{C}^n$ s.t. $w^T(\lambda I - A) = 0$: $w^T A = \lambda w^T$ (w is the left eigenvector)
 - rows are dependent
- real A can still have complex e-pairs
- A, λ real $\Rightarrow \lambda$ is associated with a real v
- conjugate (negate imaginary term of complex number[s])
- hermitian conjugate (and transpose)

10.3 Scaling interpretation

Av is simply scaled version of v (λ times); all components get magnified by the same amount

10.4 Dynamic interpretation

For $Av = \lambda v$, if $\dot{x} = Ax, x(0) = v \Rightarrow \boxed{x(t) = e^{\lambda t}v} = e^{tA}v$.

- $A^2v = \lambda^2v$
- So you just need a scalar in front of the v to calculate $x(t)$!
- An eigenvector is an initial condition $x(0)$ for which the entire trajectory is really simple.
- solution $x(t) = e^{\lambda t}v$ is a mode of $\dot{x} = Ax$ (associated with eigenvalue λ)

10.5 Invariant set

a set $S \subseteq \mathbb{R}^n$ is *invariant* under $\dot{x} = Ax$ if whenever $x(t) \in S$, then $x(\tau) \in S$ for all $\tau \geq t$ (you stay stuck within the set)

- vector field interpretation: trajectories only cut *into* S

If a single point is an invariant set, it must be in the nullspace; $S = \{x_0\} \Leftrightarrow x_0 \in \mathcal{N}(A)$, so $Ax_0 = 0 = \dot{x}$.

- line $\{tv | t \in \mathbb{R}\}$ is invariant for eigenvector v

10.6 Complex eigenvectors

- for $a \in \mathbb{C}$, complex trajectory $ae^{\lambda t}v$ satisfies $\dot{x} = Ax$, as well as *real* part

$$\begin{aligned} x(t) &= \text{Re}(ae^{\lambda t}v) \\ &= e^{\sigma t} \begin{bmatrix} v_{re} & v_{im} \end{bmatrix} \begin{bmatrix} \cos \omega t & \sin \omega t \\ -\sin \omega t & \cos \omega t \end{bmatrix} \begin{bmatrix} \alpha \\ -\beta \end{bmatrix} \end{aligned}$$

where

$$v = v_{re} + jv_{im}, \lambda = \sigma + j\omega, a = \alpha + j\beta$$

- σ gives logarithmic growth/decay factor
- ω gives angular velocity of rotation in plane
- trajectory stays in *invariant plane* $\text{span}\{v_{re}, v_{im}\}$

10.7 Dynamic interpretation: left eigenvectors

10.8 Summary:

- *right eigenvectors* are initial conditions from which resulting motion is simple (i.e., remains on line or in plane)
- *left eigenvectors* give linear functions of state that are simple, for any initial condition

10.9 Example- companion matrix

- Easy to get the characteristic polynomial
- General truth: with these matrices you can't generally tell the system behavior by just looking at it
- If you push a signal through an integrator, it gets less wiggly
- By multiplying by the left eigenvector, you've filtered out the sinusoid?

11 Lecture 13

11.1 Example: Markov chain

Probability vector $p \in \mathbb{R}^n$ that you're in each of n states: $p(t+1) = Pp(t)$. This probability evolves in time by being multiplied by state transition matrix P .

- $p_i(t) = \mathbf{Prob}(z(t) = i) \Rightarrow \sum_{i=1}^n p_i(t) = 1$
- sum of each column is 1
 - called stochastic
- i.e., $[1 \ 1 \ \cdots \ 1]$ is a left eigenvector of P with $\lambda = 1$
- so $\det(I - P) = 0$, so there's also a nonzero right eigenvector s.t. $Pv = v$
 - v can always be chosen to have non-negative elements, and can be normalized
- **Interpretation:** v is an equilibrium distribution; you don't change your *probability* distribution in time; always in v
 - if v unique, it's called the steady-state distribution of the Markov chain

11.2 Diagonalization

- v_1, \dots, v_n is LI set of eigenvectors of $A \in \mathbb{R}^{n \times n}$: $Av_i = \lambda_i v_i$
- Concatenate in matrix language:

$$A \begin{bmatrix} v_1 & \cdots & v_n \end{bmatrix} = \begin{bmatrix} v_1 & \cdots & v_n \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \lambda_2 & \\ & & \ddots \\ & & & \lambda_n \end{bmatrix}$$

or, $AT = T\Lambda$, or $T^{-1}AT = \Lambda$

- note, T is invertible, since its columns are linearly independent
- This is why, while $Av = \lambda v$ is more commonly used for a scalar eigenvalue, $\boxed{Av = v\lambda}$ is more general, as it can represent a vector of eigenvalues λ .
- so, A is diagonalizable if
 - $\exists T$ s.t. $T^{-1}AT = \Lambda$ is diagonal
 - A has a set of linearly independent eigenvectors
 - * if A not diagonalizable, it is called defective

11.3 Not all matrices diagonalizable

i.e.,

$$A = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$$

11.4 Distinct eigenvalues

fact: distinct eigenvalues in $A \Rightarrow A$ diagonalizable

- converse not true, i.e., $I \in \mathbb{R}^{7 \times 7}$

11.5 Diagonalization and left eigenvectors

rewrite $T^{-1}AT = \Lambda$ as $T^{-1}A = \Lambda T^{-1}$:

$$\begin{bmatrix} w_1^T \\ \vdots \\ w_n^T \end{bmatrix} A = \Lambda \begin{bmatrix} w_1^T \\ \vdots \\ w_n^T \end{bmatrix}$$

- remember that Λ is diagonal matrix, and multiplying by a diagonal matrix on the left is equivalent to scaling rows of the matrix
 - on the right scales the columns

Remember left/right multiplication results (whether it scales columns or rows) with 2×2 matrix multiplication:

$$\begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} x_1 & x_2 \\ y_1 & y_2 \end{bmatrix} = \begin{bmatrix} 2x_1 & 2x_2 \\ 3y_1 & 3y_2 \end{bmatrix}$$

I.e., right multiplication of diagonal matrix scales the rows.

- Take LI set of eigenvectors as columns, invert that matrix, then the rows are **left** eigenvectors
- An eigenvector is still an eigenvector after being scaled; so any can be normalized

11.6 Modal form

Take a LI set of eigenvectors from A , shove them together as columns of new matrix $T = "A \text{ is diagonalizable by } T"$

- can define new coordinates by $x = T\tilde{x}$:
- \tilde{x} is coordinates of x in the T expansion; modal (or eigenvector) expansion
 - \tilde{x} is x in terms of the eigenvectors

$$T\dot{\tilde{x}} = AT\tilde{x} \Leftrightarrow \dot{\tilde{x}} = T^{-1}AT\tilde{x} \Leftrightarrow \dot{\tilde{x}} = \Lambda\tilde{x}$$

- in new coordinate system, system is diagonal (decoupled)
- normally, with $\dot{x} = Ax$, there's a ton of cross-gains from input x_i to output y_j , where all the outputs depend on all the inputs (assuming A has only non-zero entries)
 - diagonalized system decouples it; trajectory consists of n independent modes:

$$\tilde{x}_i(t) = e^{\lambda_i t} \tilde{x}_i(0)$$

11.7 Real modal form

when eigenvalues ($\Rightarrow T$) are complex

- notes show block diagram of complex mode (note if real parts σ are removed, you get harmonic oscillator)

11.8 Diagonalization simplification

Simplifies calculation of:

- resolvent
- powers (A^k)
- exponential ($e^A = T \mathbf{diag}(e^{\lambda_1}, \dots, e^{\lambda_n}) T^{-1}$)
- So, diagonalization is largely a conceptual tool, and sometimes gives great computational advantage

11.9 Simplify for analytical functions of a matrix

11.10 Solution via diagonalization

$\dot{x} = Ax$ solution is $x(t) = e^{tA}x(0)$

- with diagonalization, solution given as

$$x(t) = \sum_{i=1}^n e^{\lambda_i t} (w_i^T x(0)) v_i$$

11.11 Interpretation

- (left eigenvectors) decompose initial state $x(0)$ into modal components $w_i^T x(0)$
- $e^{\lambda_i t}$ term propagates i th mode forward t seconds
- reconstruct state as linear combination of (right eigenvectors)

11.12 Application

Finding $x(0)$ that gives stable solution.

11.13 Stability of discrete-time systems

- powers of complex numbers s^k go to zero if $|s| < 1$
 - imaginary part tells how much of a rotation at each step you get
- $x(t+1) = Ax(t)$ is stable iff all eigenvalues of A have magnitude less than one
- spectral radius of A : $\rho(A) = \max |\lambda_i|$
 - so it is a stable system iff $\rho(A) < 1$
 - ρ gives rough growth or decay

11.14 Jordan Canonical form

- Any matrix $A \in \mathbb{R}^{n \times n}$ can be expressed in Jordan-canonical form (via ‘similarity transformation,’ for some invertible matrix T^{-1})

$$T^{-1}AT = J = \begin{bmatrix} J_1 & & \\ & \ddots & \\ & & J_q \end{bmatrix}$$

where

$$J_i = \begin{bmatrix} \lambda_i & 1 & & \\ & \lambda_i & \ddots & \\ & & \ddots & 1 \\ & & & \lambda_i \end{bmatrix} \in \mathbb{C}^{n_i \times n_i}$$

- J is ‘upper bidiagonal’
- Jordan form is unique (up to permutations of blocks- blocks might be in different places in the diagonal)
- *Almost* strictly a conceptual tool; almost never used for numerical computations
- Jordan forms are inutil if the matrix is already diagonalizable

- When you get into Jordan form, you can use a chain of integrators to represent it in block diagram form
- Jordan blocks refer to dynamics blocks that cannot be decoupled
- Jordan blocks yield:
 - repeated poles in resolvent
 - terms of form $t^p e^{t\lambda}$ in e^{tA} : solution to $\dot{x} = Ax$ is in form of polynomials, with $t^n e^{\lambda t}$ terms

11.15 Jordan block example

$$\dot{x} = \begin{bmatrix} 0 & 1 & & \\ & 0 & 1 & \\ & & 0 & 1 \\ & & & 0 & 1 \end{bmatrix} x$$

- Eigenvalues are all zero
- in the solutions you should expect $t^p e^0$ ('constant') terms: solution of form $t + t^2 + t^3$

12 Lecture 14

- Jordan canonical form: 'closest you can get to diagonalizable form, if matrix isn't diagonalizable'
- via inverse Laplace transform:

$$e^{tJ_\lambda} = e^{t\lambda}(I + tF_1 + \dots + (t^{k-1}/(k-1)!)F_{k-1})$$

$$= e^{t\lambda} \begin{bmatrix} 1 & t & \dots & t^{k-1}/(k-1)! \\ & 1 & \dots & t^{k-2}/(k-2)! \\ & & \ddots & \vdots \\ & & & 1 \end{bmatrix}$$

Jordan-blocks should be associated in your mind with polynomials times $e^{\lambda t}$

- All FIR filters have Jordan blocks

12.1 J-C application: Caley-Hamilton theorem

- if proving things about matrices, first show it for diagonalizable matrices, then with Jordan-Canonical form
- for any $n \times n$ matrix A , the powers of $A : I, A^2, A^3, \dots, A^n$ span $\mathbb{R}^{n \times n}$
 - $\mathcal{X}(A) = 0$ (characteristic polynomial; by the time you get to n , you're getting dependent matrices)
- wrong proof for evaluating characteristic polynomial $\mathcal{X}(s) = \det(sI - A)$ at A :
 - $\mathcal{X}(A) = \det(AI - A) = \det(0) = 0$
 - wrong, since definition of \mathcal{X} involves sI term, which is diagonal matrix (so A can't be plugged into it)
- correct example for

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

$$\mathcal{X} = A^2 - 5A - 2I \tag{8}$$

$$= \begin{bmatrix} 7 & 10 \\ 15 & 22 \end{bmatrix} - 5 \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} - 2I \tag{9}$$

$$= 0 \tag{10}$$

- Caley-Hamilton theorem: the square of a (square) matrix is a linear combination of the Identity matrix and itself
 - the inverse is also a linear combination of I and powers of A :
 - rewrite C-H theorem: $\mathcal{X}(A) = A^n + a_{n-1}A^{n-1} + \dots + a_0I = 0$ as
 - * $I = A(-(a_1/a_0)I - (a_2/a_0)A - \dots - (1/a_0)A^{n-1})$

$$A^{-1} = -(a_1/a_0)I - (a_2/a_0)A - \dots - (1/a_0)A^{n-1}$$

- Corollary: for every $p \in \mathbb{Z}_+$,

$$A^p \in \text{span}\{I, A, A^2, \dots, A^{n-1}\}$$

12.2 Solving huge matrix equations

- usually have a method for fast matrix vector multiplication: Az
- if you can do a quick Az calculation, you can also do a quick $A^2z, \dots, A^{n-1}z$ calculation
 - using C-H thm, you can put scalars in front of them to get A^{-1}

12.3 LTI system inputs & outputs

LTI systems with input and outputs have the form $\dot{x} = Ax + Bu, y = Cx + Du$; Ax is the *drift* term, Bu is the *input* term. For

12.4 Interpretations

- state derivative \dot{x} is the sum of the autonomous term Ax and one term per input, $b_i u_i$
 - if columns of B are independent, each input u_i gives another degree of freedom for \dot{x}
 - column i of B represents how much input i affects \dot{x}
 - row i of B represents how much input i affects the i th component of \dot{x}
- Good block diagram in notes, pg. 13-5

12.5 Transfer matrix

- Solution to $\dot{x} = Ax + Bu$ via Laplace transform is $x(t) = e^{tA}x(0) + \int_0^t e^{(t-\tau)A}Bu(\tau)d\tau$
 - $e^{tA}x(0)$ is the (unforced or) autonomous response
 - $e^{tA}B$ is the ‘input-to-state impulse matrix’
 - * integral term is convolution
 - $(sI - A)^{-1}B$ is input-to-state *transfer matrix*; the resolvent $\times B$
- Readout equation, $y = Cx + Du$ solution (via Laplacian)
 - $y(t) = Ce^{tA}x(0) + \int_0^t Ce^{(t-\tau)A}Bu(\tau)d\tau + Du(t)$
 - transfer function: $H(s) = C(sI - A)^{-1}B + D$
 - impulse matrix/impulse response: $h(t) = Ce^{tA}B + D\delta(t)$
 - zero initial condition gives $Y(s) = H(s)U(s), y = h * u$
 - * H_{ij} is transfer function from input u_j to output y_i

- Mnemonic for convolution intuition $y(t) = \int_0^t h(t-\tau)u(\tau)d\tau$:
 - output is ‘linear combination’ or ‘mixture’ of input in the past $u(\tau)$, weighted by how many seconds ago something happened
 - τ is ‘seconds ago’
 - example: if $h(7)$ is really big \rightarrow output depends a *lot* on the input 7 seconds ago
 - * remember, the h term has been delayed and flipped: $h(-(\tau-t))$
 - if $h(t)$ decays, the system has *fading memory*

12.6 Impulse matrix

- $h(t) = Ce^{tA}B + D\delta(t), x(0) = 0, y = h * u$

$$y_i(t) = \sum_{j=1}^m \int_0^t h_{ij}(t-\tau)u_j(\tau)d\tau$$

- for rain/river metaphor, where rain is input and river level is output, this matrix is summed over the input, with different regions of rainfall u_j and different locations of the river where rainfall is measured, $y_i(t)$

12.7 Step response

$$s(t) = \int_0^t h(\tau)d\tau$$

12.8 Spring dashpot example

12.9 RC circuit example

13 Lecture 15

13.1 DC or static gain matrix

- transfer matrix at $s = 0$ is $H(0) = -CA^{-1}B + D \in \mathbb{R}^{m \times p}$
- DC transfer matrix describes system under static conditions: x, u, y constant
 - $\Rightarrow y = H(0)u$
 - if system is stable, $H(0) = \int_0^\infty h(t)dt = \lim_{t \rightarrow \infty} s(t)$

13.2 Discretization with piecewise constant inputs

13.3 Causality

13.4 Idea of state

- $x(t)$ is state of system at time t :
 - future output depends only on *current state* and *future output*
 - * sufficient statistic of what happened in the past to be able to predict the future
 - bridge between past inputs and future outputs
 - * ‘past and future are conditionally independent, given the state’ (-machine learning)

13.5 Change of coordinates

- start with LDS $\dot{x} = Ax + Bu, y = Cx + Du$
- change coordinates in \mathbb{R}^n to \tilde{x} with $x = T\tilde{x}$

$$\dot{\tilde{x}} = T^{-1}\dot{x} = T^{-1}(Ax + Bu) = T^{-1}AT\tilde{x} + T^{-1}Bu$$

- LDS can be expressed as

$$\dot{\tilde{x}} = \tilde{A}\tilde{x} + \tilde{B}u, y = \tilde{C}\tilde{x} + \tilde{D}u$$

where

$$\tilde{A} = T^{-1}AT, \tilde{B} = T^{-1}B, \tilde{C} = CT, \tilde{D} = D$$

- You might want to do this to, i.e., get an A with a bunch of zeros, since the system might be much easier to implement

13.6 Standard forms for LDS

Can change coordinates to put A in various forms (diagonal, real modal, Jordan ...)

13.7 Discrete-time systems

(Another block diagram in notes)

13.8 Z-transform

13.9 Discrete-time transfer function

- re-derivation of stuff we know, using the frequency domain

14 Last Section, Singular Value Decomposition

14.1 Eigenvalues of symmetric matrices

Given $A \in \mathbb{R}^{n \times n} : A = A^T$

- eigenvalues always real
- $Av = \lambda v, v \neq 0, v \in \mathbb{C}^n$; remember conjugate transpose, $\bar{v}^T = v^H = v^{*T}$, or, in MATLAB,

v'

and remember, $\bar{a}a = |a|^2$

$$\Rightarrow v^H Av = v^H (Av) = \lambda v^H v = \lambda \sum_{i=1}^n |v_i|^2$$

$$\dots = \bar{\lambda} \sum_{i=1}^n |v_i|^2$$

so $\lambda = \bar{\lambda} \Rightarrow \lambda \in \mathbb{R}$

14.2 Eigenvectors of symmetric matrices

fact: There is a set of orthonormal eigenvectors of A , $q_1, \dots, q_n : Aq_i = \lambda_i q_i, q_i^T q_j = \delta_{ij}$

- In matrix form, \exists orthogonal $Q : Q^{-1}AQ = Q^T AQ = \Lambda$
- or, rewrite *dyadic expansion* of A :

$$A = Q\Lambda Q^T = \sum_{i=1}^n \lambda_i (q_i q_i^T)$$

- *engineering etiquette aside*: “The eigenvectors of a symmetric matrix are orthonormal” makes no sense; any matrix has zillions of eigenvectors
 - but rather, “you can choose the eigenvectors of a symmetric matrix to be orthonormal”

14.3 Interpretations

Given $A = Q\Lambda Q^T$. Remember, $Q^{-1} = Q^T$. (Notes have block diagram.) This means, to multiply a vector x by A , first multiply it by Q^T , then Λ , then Q

- first operation result $Q^T x = Q^{-1}x$ ‘resolves x in the q_i coordinates’
- next: simple to multiply by diagonal matrix Λ (simply scaling the matrix)
- last: multiplying by Q ‘reconstitutes the output’
- multiplying by Q^T, Λ, Q represents a ‘coding,’ scaling, reconstruction operation
- Application: in JPEG, do DCT transformation, quantize in middle, then decode the image in last step
- Geometrically:
 - rotate by Q^T
 - dilation by Λ
 - rotate back by Q

Decomposition (review)

$$A = \sum_{i=1}^n \lambda_i q_i q_i^T$$

expresses A as linear combination of 1-dimensional projections.

If a matrix is real and symmetric, (and the eigenvalues of a matrix are distinct,) then any set of associated eigenvectors is orthogonal (and can be normalized).

$$Av_i = \lambda_i v_i, \|v_i\| = 1$$

Then,

$$v_i^T (Av_j) = \lambda_j v_i^T v_j = (Av_i)^T v_j = \lambda_i v_i^T v_j$$

so $(\lambda_i - \lambda_j)v_i^T v_j = 0$ for $i \neq j, \lambda_i \neq \lambda_j \Rightarrow v_i^T v_j = 0$ So distinct eigenvectors of a symmetric real matrix are ‘orthogonal’

14.4 Example: RC circuit

15 Lecture 16

15.1 Quadratic forms

Previously, we did mappings of form $y = Ax$, or $\dot{x} = Ax$. We're no longer doing problems with a linear structure.

- New form is *quadratic form*

$$f(x) = x^T A x = \sum_{i,j=1}^n A_{ij} x_i x_j$$

- $(x^T A x)^T = x^T A^T x$
- $x^T A x = x^T ((A + A^T)/2) x$ is another form of quadratic form; the middle part is the average of itself (member of A) and its associated transposed self
- weird etiquette: A should be symmetric if you're dealing with quadratic forms
 - often may want to symmetrize it to make it safe
 - i.e., this quadratic form

$$x^T \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} x = x_1 x_2$$

$$x^T \begin{bmatrix} 0 & 3/2 \\ -1/2 & 0 \end{bmatrix} x = x_1 x_2$$

But the canonical form would have 1/2 in the 2nd and 3rd entry (row-wise).

- If A is diagonal in LDS, $y = Ax$, it means there is no cross-coupling. If A is diagonal in a quadratic form, it gives you a weighted sum of squares.
- **uniqueness:** if $x^T A x = x^T B x$ for all $x \in \mathbb{R}^n$ and $A = A^T, B = B^T$, then $A = B$

15.2 Examples of quadratic forms

- $\|Bx\|^2 = x^T (B^T B) x$; $B^T B$ is already symmetric
- $\sum_{i=1}^{n-1} (x_{i+1} - x_i)^2$. This measures 'wiggleness' of signal; sum of squares of difference from one to the next signal
 - to get in quadratic form, multiply out: $x_{i+1}^2 + x_i^2 - 2x_i x_{i+1}$:

$$\begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} 1 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

- Zeros in corners of matrix, since there is no product between x_1, x_3
- $\|Fx\|^2 - \|Gx\|^2$
 - $x^T (F^T F - G^T G) x$
- quadratic surface: level set: $\{x | f(x) = a\}$
- quadratic region: $\{x | f(x) \leq a\}$
 - unit ball: $\{x | x^T x \leq 1\}$
 - unit sphere: $\{x | x^T x = 1\}$

15.3 Inequalities for quadratic forms

If $A = A^T$, it is diagonalizable: $A = Q\Lambda Q^T$ with sorted eigenvalues in Λ

$$x^T A x = x^T Q \Lambda Q^T x \quad (11)$$

$$= (Q^T x)^T \Lambda (Q^T x) \quad (12)$$

$$= \sum_{i=1}^n \lambda_i (q_i^T x)^2 \quad (13)$$

$$\leq \lambda_1 \sum_{i=1}^n (q_i^T x)^2 \quad (14)$$

$$= \lambda_1 \|x\|^2 \quad (15)$$

$$(16)$$

- So, how big can the quadratic form be?

$$- x^T A x \leq \lambda_1 x^T x$$

- Another important inequality

$$\lambda_n x^T x \leq x^T A x \leq \lambda_1 x^T x$$

- λ_1 sometimes called λ_{\max} , λ_n called λ_{\min}

- for $\boxed{A = A^T} \in \mathbb{R}^{n \times n}$, A is *positive semidefinite* if $x^T A x \geq 0$ for all x

– notation: $A \geq 0$

– iff $\lambda_{\min}(A) \geq 0$: all eigenvalues nonnegative

– A is *positive definite* if $x^T A x > 0$ for all $x \neq 0$

- $A \geq 0 \Rightarrow A_{ij} \geq 0$

- if A not positive or negative [semi]definite, it is *indefinite*

- matrices **can be incomparable**:

– $A \not\geq B$ and $B \not\geq A$, where it will depend on the x (depend on the direction you're looking at it from)

- $A \leq B$ statement requires that both be square matrices, and probably symmetric (depending on social norms)

– if not symmetric, it means $\frac{A+A^T}{2} \leq \frac{B+B^T}{2}$

15.4 Ellipsoids

- if $A^T > 0$, the set $\mathcal{E} = \{x | x^T A x \leq 1\}$ is an *ellipsoid* in \mathbb{R}^n , centered at 0

- semi-axes given by $s_i = \lambda_i^{-1/2} q_i$

– eigenvectors determine directions of semi-axes

– eigenvalues determine lengths of semi-axes

– $\sqrt{\lambda_{\max}/\lambda_{\min}}$ gives maximum eccentricity

- eigenvalue intuition in \mathbb{R}^3

– large pancake shape \Rightarrow one large eigenvalue, two slightly smaller eigenvalues

– long cigar shape \Rightarrow two large eigenvalues, one small eigenvalue

15.5 Gain of matrix in a direction

For $A \in \mathbb{R}^{m \times n}$ (not necessarily symmetric anymore), for $x \in \mathbb{R}^n$, $\|Ax\|/\|x\|$ gives *amplification factor* or *gain* of A in direction x , and will vary with different x directions

- Questions we will ask

- maximum gain of A

- * $\max_{x \neq 0} \frac{\|Ax\|}{\|x\|}$

- * *matrix norm* or *spectral norm*, and denoted $\|A\|$ (overloading)

- * $\max_{x \neq 0} \frac{\|Ax\|}{\|x\|} = \max_{x \neq 0} \frac{x^T A^T A x}{\|x\|^2} = \lambda_{\max}(A^T A)$

- * so, $\|A\| = \sqrt{\lambda_{\max}(A^T A)}$

- minimum gain of A

- * $\min_{x \neq 0} \|Ax\|/\|x\| = \sqrt{\lambda_{\min}(A^T A)}$

- how does gain of A vary with direction

- Will give us a *quantitative way* to talk about nullspace

15.6 Example

15.7 Properties of matrix norm

- triangle inequality
- scaling
- definiteness: $\|A\| = 0 \Leftrightarrow A = 0$
- norm of a product inequality

16 Lecture 17

16.1 SVD

Singular value decomposition, for *any* matrix A .

$$A = U \Sigma V^T$$

where

- $A \in \mathbb{R}^{m \times n}$, $\mathbf{Rank}(A) = r$
- $U \in \mathbb{R}^{m \times r}$, $U^T U = I$
- $V \in \mathbb{R}^{n \times r}$, $V^T V = I$
- $\Sigma = \mathbf{diag}(\sigma_1, \dots, \sigma_r)$, where $\sigma_1 \geq \dots \geq \sigma_r > 0$

16.2 Dyadic expansion: (U and V have orthonormal columns)

$$A = U\Sigma V^T = \sum_{i=1}^r \sigma_i u_i v_i^T$$

- σ_i are the nonzero *singular values* of A
- v_i are the *right* or *input singular vectors* of A
- u_i are the *left* or *output singular vectors* of A

$$A^T A = (U\Sigma V^T)^T (U\Sigma V^T) = V\Sigma^2 V^T$$

- v_i are eigenvectors of $A^T A$ corresponding to nonzero eigenvalues
- $\sigma_i = \sqrt{\lambda_i(A^T A)}$ (and $\lambda_i(A^T A) = 0$ for $i > r$)
- $\|A\| = \sigma_1$

$$AA^T = (U\Sigma V^T)(U\Sigma V^T)^T = U\Sigma^2 U^T$$

- u_i are eigenvectors of AA^T corresponding to nonzero eigenvalues
- $\sigma_i = \sqrt{\lambda_i(AA^T)}$ (and $\lambda_i(AA^T) = 0$ for $i > r$)
- u_1, \dots, u_r are orthonormal basis for $\text{range}(A)$
- v_1, \dots, v_r are orthonormal basis for $\mathcal{N}(A)^\perp$

16.3 Interpretations

$$A = U\Sigma V^T$$

linear mapping $y = Ax$ can be decomposed as

- compute coefficients of x along input directions v_1, \dots, v_r
- scale coefficients by σ_i
- reconstitute along output directions u_1, \dots, u_r
- u_1 is highest gain output direction
- v_1 is most sensitive (highest gain) input direction
- $Av_1 = \sigma_1 u_1$

SVD gives clearer picture of gain as function of input/output directions.

16.4 General pseudo-inverse

- $A \neq 0 \Rightarrow A = U\Sigma V^T$

General pseudo-inverse: $A^\dagger = V\Sigma^{-1}U^T$ is the pseudo-inverse of A .

If A is skinny and full rank,

$$A^\dagger = (A^T A)^{-1} A^T$$

If A is fat and full rank,

$$A^\dagger = A^T (AA^T)^{-1}$$

In the general case, (if A is not full-rank)

$$X_{ls} = \{z : \|Az - y\| = \min_w \|Aw - y\|\}$$

is the set of least squares solutions.

$x_{pinv} = A^\dagger y \in X_{ls}$ has minimum norm on X_{ls} , so this is the minimum-norm, least squares solution

16.5 Pseudo-inverse via regularization

Minimizer of $\|Ax - y\|^2 + \mu\|x\|^2$ is given as

$$x_\mu = (A^T A + \mu I)^{-1} A^T y$$

17 Lecture 18

17.1 Sensitivity of linear equations to data error

To perturbations $\delta x, \delta y$. Looking at *relative error* gives us

$$\frac{\|\delta x\|}{\|x\|} \leq \|A\| \|A^{-1}\| \frac{\|\delta y\|}{\|y\|}$$

$$\kappa(A) = \text{cond}(A) = \|A\| \|A^{-1}\| = \sigma_{\max}(A) / \sigma_{\min}(A)$$

So, the relative error in solution $x \leq$ condition number \cdot relative error in data y . In terms of # of bits of guaranteed accuracy, ‘# of bits accuracy in solution \approx # bits accuracy in data $-\log_2 \kappa$.’

17.2 Low rank applications

If you want to decompose A into $A = BC$, you can remember QR factorization (does skinny-fat factorization), but we’d also want an A that is approximated by such a factorization. Suppose $A \in \mathbb{R}^{m \times n}$, $\mathbf{Rank}(A) = r$ with SVD $A = U \Sigma V^T = \sum_{i=1}^r \sigma_i u_i v_i^T$, and we want $\hat{A} : \mathbf{Rank}(\hat{A}) \leq p < r : \|A - \hat{A}\|$ is minimized.

The optimal rank p approximator of A is

$$\hat{A} = \sum_{i=1}^p \sigma_i u_i v_i^T$$

- $\|A - \hat{A}\| = \|\sum_{i=p+1}^r \sigma_i u_i v_i^T\| = \sigma_{p+1}$
- SVD dyads $u_i v_i^T$ are ranked in order of ‘importance’; take p to get rank p approximant
- keep in mind Frobenius norm $\|A - \hat{A}\|_F = \|A(i) - \hat{A}(i)\| = \sqrt{\sum_{i,j} (A_{ij} - \hat{A}_{ij})^2}$
- If you have a set of 100 \mathbb{R}^{10} prices that has dimension=3, it means there are 3 underlying factors; this would give you ‘stunning’ predictive power
- You would *almost* (i.e., practically) have a dimension 3 matrix if the SVD Σ had 3 singular values that were significantly larger than the rest

17.3 State transfer

17.4 Reachability

Consider state transfer from $x(0) = 0$ to $x(t)$

- $x(t)$ is reachable in t seconds
- define $\mathcal{R}_t \subseteq \mathbb{R}^n$ as the set of points reachable in t seconds for CT system $\dot{x} = Ax + Bu$,

$$\mathcal{R}_t = \left\{ \int_0^t e^{(t-\tau)A} B u(\tau) d\tau \mid u : [0, t] \rightarrow \mathbb{R}^m \right\}$$

and for DT system $x(t+1) = Ax(t) + Bu(t)$,

$$\mathcal{R}_t = \left\{ \sum_{\tau=0}^{t-1} A^{t-1-\tau} B u(\tau) \mid u(t) \in \mathbb{R}^m \right\}$$

If a state is reachable in 1 s, it’s reachable in 2 s (just leave it at zero for the first second) if you’re starting from zero initial state (in both CT and DT systems).

17.5 Reachability for discrete-time LDS

DT system $x(t+1) = Ax(t) + Bu(t), x(t) \in \mathbb{R}^n$

$$x(t) = \mathcal{C}_t \begin{bmatrix} u(t-1) \\ \vdots \\ u(0) \end{bmatrix}$$

where $\mathcal{C}_t = [B \quad Ab \quad \dots \quad A^{t-1}B]$. So reachable set at t is $\mathcal{R}_t = \text{range}(\mathcal{C}_t)$

If we want to see if a state is reachable in 1 step, see if it's in the range of B .

If we want to see if it's reachable in 2 steps, see if it's in the range of B, AB .

Caley-Hamilton theorem lets us express each A^k for $k \geq n$ as linear combination of A^0, \dots, A^{n-1} .

18 Lecture 19

But, you can't increase your reach after $t \geq n$ seconds (or epochs). I.e.,

$$\mathcal{R}_t = \begin{cases} \text{range}(\mathcal{C}_t) & t < n \\ \text{range}(\mathcal{C}) & t \geq n \end{cases}$$

where $\mathcal{C} = \mathcal{C}_n$ is the *controllability matrix*.

18.1 General state transfer

$$x(t_f) = A^{t_f - t_i} x(t_i) + \mathcal{C}_{t_f - t_i} \begin{bmatrix} u(t_f - 1) \\ \vdots \\ u(t_i) \end{bmatrix}$$

- First term on rhs is what would happen without any input (the drift term)

so we can transfer $x(t_i)$ to $x(t_f) = x_{des}$ iff

$$x_{des} - A^{t_f - t_i} x(t_i) \in \mathbb{R}_{t_f - t_i}$$

- we're not hoping x_{des} is reachable, we want it minus the drift
 - So it is possible to be able to reach a state in 4 steps, but not in 5
- This reduces to the reachability problem
- if system is controllable, any state transfer can be achieved in $\leq n$ steps
- important special case: driving state to zero

18.2 Minimum time control problem

Given

$$x(t+1) = Ax(t) + Bu(t), \quad x(0) = x_0$$

how do you minimize t ? Check if $A^t x_0 \in \mathcal{R}([B \quad \dots \quad A^{t-1}B])$

18.3 Least-norm input for reachability

assume system is reachable, $\text{Rank}(\mathcal{C}_t) = n$ to steer $x(0) = 0$ to $x(t) = x_{des}$, inputs $u(0), \dots, u(t-1)$ to satisfy

$$x_{des} = \mathcal{C}_t \begin{bmatrix} u(t-1) \\ \vdots \\ u(0) \end{bmatrix}$$

among all u that steer $x(0) = 0$ to $x(t) = x_{des}$, the one that minimizes

$$\sum_{\tau=0}^{t-1} \|u(\tau)\|^2$$

The least-square solution for the input is given by

$$\begin{bmatrix} u_{ln}(t-1) \\ \vdots \\ u_{ln}(0) \end{bmatrix} = \mathcal{C}_t^T (\mathcal{C}_t \mathcal{C}_t^T)^{-1} x_{des}$$

18.4 Stability

18.5 Continuous time reachability

consider $\dot{x} = Ax + Bu$ with $x(t) \in \mathbb{R}^n$
reachable set at time t is

$$\mathcal{R}_t = \left\{ \int_0^t e^{(t-\tau)A} Bu(\tau) d\tau \mid u : [0, t] \rightarrow \mathbb{R}^m \right\}$$

for $t > 0$, $\mathcal{R}_t = \mathcal{R} = \text{range}(\mathcal{C})$ where

$$\mathcal{C} = [B \ AB \ \cdots \ A^{n-1}B]$$

is the controllability matrix of (A, B)

- According to the model, you can hit anything infinitely fast

18.6 Impulsive inputs

18.7 Example

18.8 Least-norm input for reachability

Assume that $\dot{x} = Ax + Bu$ is reachable. We want u that steers $x(0) = 0$ to $x(t) = x_{des}$ and minimizes

$$\int_0^t \|u(\tau)\|^2 d\tau$$

let's discretize system with interval $h = t/N$ (eventually let $N \rightarrow \infty$). So u is piecewise constant:

$$u(\tau) = u_d(k) \text{ for } kh \leq \tau < (k+1)h, \quad k = 0, \dots, N-1$$

... Get long expression for least norm solution for u

19 Lecture 20

19.1 Observability and state estimation

Consider DT system

$$x(t+1) = Ax(t) + Bu(t) + w(t), \quad y(t) = Cx(t) + Du(t) + v(t)$$

- w is state disturbance/noise
- v is sensor noise/error
- state estimation problem: estimate $x(s)$ from

$$u(0), \dots, u(t-1), \quad y(0), \dots, y(t-1)$$

- $s = 0$: estimate initial state
- $s = t - 1$: estimate current state
- $s = t$: estimate/predict next state
- estimate $\tilde{x}(s)$, called *observer* or *state estimator*

19.2 Noiseless case

Find $x(0)$, with no state or measurement noise

$$x(t+1) = Ax(t) + Bu(t), \quad y(t) = Cx(t) + Du(t)$$

then we have

$$\begin{bmatrix} y(0) \\ \vdots \\ y(t-1) \end{bmatrix} = \mathcal{O}_t x(0) + \mathcal{T}_t \begin{bmatrix} u(0) \\ \vdots \\ u(t-1) \end{bmatrix}$$

where

$$\mathcal{O}_t = \begin{bmatrix} C \\ CA \\ \vdots \\ CA^{t-1} \end{bmatrix} \quad \mathcal{T}_t = \begin{bmatrix} D & 0 & \dots & \dots \\ CB & D & 0 & \dots \\ \vdots & \vdots & \vdots & \vdots \\ CA^{t-2}B & CA^{t-3}B & \dots & CB & D \end{bmatrix}$$

- \mathcal{O}_t maps initial state into resulting output over $[0, t-1]$
- \mathcal{T}_t maps input to output over $[0, t-1]$

19.3 Observability matrix

19.4 Least-squares observers

19.5 Parting thoughts

19.5.1 Linear algebra

19.5.2 Levels of understanding

- High school: 17 variables, 17 equations \rightarrow usually has unique solution
 - 80 vars, 60 eqns \rightarrow probably 20 extra degrees of freedom
- Platonic view ('math')
 - singular, rank, range, nullspace, Jordan form, controllability
 - stuff is true or false
- Quantitative
 - based on least-squares, SVD
 - gives numerical measures for ideas like singularity, rank, etc
 - interpretation depends on (practical) context
 - very useful in practice

20 Appendix

Some special things to remember.

20.1 Inverse, transpose properties

- $(AB)^{-1} = B^{-1}A^{-1}$
- $(A^{-1})^T = (A^T)^{-1} = A^{-T}$

20.2 Invertibility implications

For an n -by- n matrix A

Invertible	mnemonic
$ A \neq 0$	$ A = 0 \Rightarrow$ you can't compute the inverse - (remember base case 2×2 matrix inverse involves $1/ A $ term)
non-singular	singular \Rightarrow the matrix sends a nontrivial subspace to the singular subspace, $\{0\}$
A is full rank	linearly independent columns (invertibility \Rightarrow 1-to-1/injective)
$\mathcal{N}(A) = \{0\}$	linearly independent columns
$\mathcal{R}(A) = \mathbb{R}^n$	linearly independent columns
$Ax = b$ has unique solution for every b	- no more than one solution (can't add members of $\mathcal{N}(A)$ for multiple b) - one solution, since $\mathcal{R}(A) = \mathbb{R}^n$; everything reachable/surjective - one solution found using the unique inverse of A
$\text{rref}(A) = I_n$	
A is a product of elementary matrices	

20.3 Tips

- to sum up all the elements in matrix A , multiply by one vectors: $\sum_{i,j} (A)_{i,j} = \mathbf{1}^T A \mathbf{1}$
- each row i in Z is a linear combination of rows i, \dots, n in $Y : Z = UY$, where U is upper triangular: $U_{ij} = 0$ for $i > j$
 - more in 2.24, hw2
- $E_{i,j} = e_i e_j^T \in \mathbb{R}^{n \times n}$, $i, j = 1, \dots, n$ is $n \times n$ matrix with a 1 in i, j th entry, and zero elsewhere
 - has n^2 dimensions
- When you do a polynomial of a similarity matrix, you can pull out the outer matrices: $\mathcal{X}(T\Lambda T^{-1}) = T\mathcal{X}(\Lambda)T^{-1}$
- If $AB = 0$, and each is full rank, (neither is zero), then $\mathcal{R}(B) \subseteq \mathcal{N}(A)$
- $A^k = 0, k \geq 1 \Rightarrow$ eigenvalues of A are zero. If A is diagonalizable, then A^k can be expressed as $T\Lambda^k T^{-1}$, where Λ is diagonal, consisting of the eigenvalues multiplied by the k th power, since T is non-singular.
 - **Rank**(A) = $r \Rightarrow A \in \mathbb{R}^{n \times n}$ has exactly r non-zero eigenvalues
- Every eigenvalue of AB is an eigenvalue of BA , if $A, B \in \mathbb{R}^{n \times n}$: Given λ is eigenvalue of $AB : ABv = \lambda v, BA(Bv) = B(AB)v = \lambda Bv$. Thus Bv is an eigenvector of BA associated with λ
- To check whether $\boxed{x \in \mathcal{R}(A)}$, check whether $\boxed{\text{rank}([A \ x]) = \text{rank}(A)}$
- Least squares approximation that minimizes x in $Ax = y, x_{ls} = (A^T A)^{-1} A^T y$ is given in MATLAB by `xls=A\y`;
 - only works if A is skinny or square

- compute using QR (economy) factorization with

```
[Q,R]=qr(A,0); % compute economy QR decomposition
xls=R\ (Q'*y);
```

- Least norm solution, $x_{ln} = A^T(AA^T)^{-1}y$, can be computed by `xln=A'*inv(A*A')*y`;

- A cannot be full rank or skinny

- via QR factorization:

```
[Q,R]=qr(A',0);
xln=Q*(R'\y);
```

* for more, see here

20.4 Unrelated stuff

20.4.1 Series

$$S = \sum_{k=0}^N a^k = a^0 + a^1 + \dots + a^N \quad (17)$$

$$aS = \sum_{k=0}^N a^k = a^1 + a^1 + \dots + a^{N+1} \quad (18)$$

$$aS - S = -1 + a^{N+1} \quad (19)$$

$$S = \frac{a^{N+1} - 1}{a - 1} \quad (20)$$

20.4.2 Taylor/Mclauren

Intuition for approximating a function at a point (say at $x = 0$): assuming we know $f(0), f'(0), f''(0), f'''(0)$, put into polynomials

$$p_0(x) = f(0) \quad (21)$$

$$p_1(x) = f(0) + f'(0)x \quad (22)$$

$$p_2(x) = f(0) + f'(0)x + \frac{1}{2}f''(0)x^2 \quad (23)$$

$$p_3(x) = f(0) + f'(0)x + \frac{1}{2}f''(0)x^2 + \frac{1}{6}f'''(0)x^3 \quad (24)$$

$$(25)$$

Our polynomial $p_n(x)$ will both have the same value at $x = 0$, as well as the same n th derivatives when evaluated at 0. I.e.,

$$p'_2(x) = f'(0) + f''(0)x \quad (26)$$

$$p'_2(0) = f'(0) \quad (27)$$

$$(28)$$

In general, the n th order McLauren approximation polynomial of $f(x)$ at point x_1 will be given as

$$f(x) \approx f(x_1) + f'(x_1)x + \frac{1}{2}f''(x_1)x^2 + \frac{1}{2 \cdot 3}f^{(3)}(x_1)x^3 + \dots + \frac{1}{n!}f^{(n)}(x_1)x^n$$

In general, the n th order Taylor approximation polynomial of $f(x)$ at point c will be given as

$$f(x) \approx f(c) + f'(c)(x - c) + \frac{1}{2}f''(c)(x - c)^2 + \dots + \frac{1}{n!}f^{(n)}(c)(x - c)^n$$

So when you evaluate derivatives at $x = c$, the higher order terms will still drop to zero, leaving you with the exact derivative value.

- Cosine MaLauren (around 0)

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

- Sine MaLauren (around 0)

$$\sin(x) = x - \frac{x^3}{3!} + \frac{x^5}{5!} - \frac{x^7}{7!} + \dots$$

- e MaLauren (around 0)

$$e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots$$

21 Homework assignments

Homework 1	Lecture 4	2.1–2.4, 2.6, 2.9, 2.12, +
Homework 2	Lecture 6	3.2, 3.3, 3.10, 3.11, 3.16, 3.17, +
Homework 3	Lecture 8	2.17, 3.13, 4.1–4.3, 5.1, 6.9, +
Homework 4	Lecture 10	5.2, 6.2, 6.5, 6.12, 6.14, 6.26, 7.3, 8.2
Homework 5	Lecture 13	10.2, 10.3, 10.4, +
Homework 6	Lecture 14	9.9, 10.5, 10.6, 10.8, 10.14, 11.3, and 11.6a
Homework 7	Lecture 16	10.9, 10.11, 10.19, 11.13, 12.1, 13.1, +
Homework 8	Lecture 18	13.17, 14.2, 14.3, 14.4, 14.6, 14.8, 14.9, 14.11, 14.13, 14.21, 14.33, +
Homework 9	Lecture 20	14.16, 14.26, 15.2, 15.3, 15.6, 15.8, 15.10, and 15.11